

NAME : SRISTI BHOWMICK , COLLEGE ROLL : BTECH/CSE/23/095 , EXAM ROLL : 11600223120

ASSIGNMENT 11

Build scaling plans in AWS that balance the load on different EC2 instances.

STEP 1: GO TO EC2 and click on launch templates -> launch templates and give the name and version type -> Give information and launch.

Ask Amazon Q



Asia Pacific (Mumbai)

Sristi Bhowmick (0217-5467-6955)

Sristi Bhowmick

EC2 launch templates

Streamline, simplify and standardize instance launches

Use launch templates to automate instance launches, simplify permission policies, and enforce best practices across your organization. Save launch parameters in a template that can be used for on-demand launches and with managed services, including EC2 Auto Scaling and EC2 Fleet. Easily update your launch parameters by creating a new launch template version.

New launch template
[Create launch template](#)

[EC2](#) > [Launch templates](#) > Create launch template

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - *required*

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description


Max 255 chars

Auto Scaling guidance | [Info](#)
Select this if you intend to use this template with EC2 Auto Scaling

☒ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

► **Template tags**

► **Source template**



[Browse more AMIs](#)
Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Ubuntu Server 24.04 LTS (HVM), SSD Volume Type
ami-05d2d839d4f73aafb (64-bit (x86)) / ami-0f6c99577b93a642b (64-bit (Arm))
Virtualization: hvm ENA enabled: true Root device type: ebs

Free tier eligible

Description

Ubuntu Server 24.04 LTS (HVM),EBS General Purpose (SSD) Volume Type. Support available from Canonical (<http://www.ubuntu.com/cloud/services>).

Canonical, Ubuntu, 24.04, amd64 noble image

Architecture

64-bit (x86)

AMI ID

ami-05d2d839d4f73aafb

Publish Date

2026-03-14

Username

ubuntu

Verified provider

▼ Instance type [Info](#) | [Get advice](#)

[Advanced](#)

Instance type

t3.micro

Free tier eligible

Family: t3 2 vCPU 1 GiB Memory Current generation: true
On-Demand Linux base pricing: 0.0112 USD per Hour On-Demand SUSE base pricing: 0.0112 USD per Hour
On-Demand Windows base pricing: 0.0204 USD per Hour
On-Demand Ubuntu Pro base pricing: 0.0147 USD per Hour On-Demand RHEL base pricing: 0.04 USD per Hour

☐ All generations

[Compare instance types](#)

Additional costs apply for AMIs with pre-installed software

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name

key2

[Create new key pair](#)

Security groups [Info](#)

Select security groups

[Compare security group rules](#)

SristiSecurity sg-0629a8fa4d7c72b46

VPC: vpc-0a874208a064ece04

For v2 requests, you must include a session token in all instance metadata requests. Applications or agents that use V1 for instance metadata access will break.

Metadata response hop limit [Info](#)

2

Allow tags in metadata [Info](#)

Don't include in launch template

User data - optional [Info](#)

Upload a file with your user data or enter it in the field.

[Choose file](#)

```
#!/bin/bash
apt-get update
apt-get install -y nginx
systemctl start nginx
systemctl enable nginx
apt-get install -y git
curl -SL https://deb.nodesource.com/setup_16.x | sudo -E bash -
apt-get install -y nodejs
git clone https://github.com/bhowmicksristi29-beep/Repository2.git
cd Repository2
npm install
node index.js
```

☐ User data has already been base64 encoded

▼ Summary

Software Image (AMI)
Canonical, Ubuntu, 24.04, amd64...[read more](#)
ami-05d2d839d4f73aafb

Virtual server type (instance type)
t3.micro

Firewall (security group)
SristiSecurity

Storage (volumes)
1 volume(s) - 8 GiB

Cancel

Create launch template



Success

Successfully created [SristiTemp\(lt-0283d209c61392956\)](#).

aws

Search

[Alt+S]

Asia Pacific

Sristi Bhowmick (0217-5467-6955)

Sristi Bhowmick

EC2

>

Auto Scaling groups

Capacity Manager

Images

AMI Catalog

Elastic Block Store

Network & Security

Load Balancing

Auto Scaling

Settings

Amazon EC2 Auto Scaling

helps maintain the availability of your applications

Auto Scaling groups are collections of Amazon EC2 instances that enable automatic scaling and fleet management features. These features help you maintain the health and availability of your applications.

Create Auto Scaling group

Get started with EC2 Auto Scaling by creating an Auto Scaling group.

Create Auto Scaling group

How it works

Auto Scaling group

aws

Search

[Alt+S]

Asia Pacific

Sristi Bhowmick (0217-5467-6955)

Sristi Bhowmick

EC2

>

Auto Scaling groups

>

Create Auto Scaling group

Step 1

Choose launch template

Step 2

Choose instance launch options

Step 3 - optional

Integrate with other services

Step 4 - optional

Configure group size and scaling

Step 5 - optional

Add notifications

Step 6 - optional

Add tags

Step 7

Review

Choose launch template

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.

Name

Auto Scaling group name

Enter a name to identify the group.

SristiAuto

Must be unique to this account in the current Region and no more than 255 characters.

Launch template

For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template

Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

SristiTemp

Create a launch template

Version

Latest (1)

Create a launch template version

Description

SristiTemp

Launch template

SristiTemp

Instance type

t3.micro

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-0a874208a064ece04
172.31.0.0/16 Default



[Create a VPC](#)

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets



aps1-az1 (ap-south-1a) | subnet-0e71abe105a1ac286
172.31.32.0/20 Default

aps1-az2 (ap-south-1c) | subnet-0e38d5355eecbfc1a
172.31.16.0/20 Default

aps1-az3 (ap-south-1b) | subnet-0f94e0905aad52965
172.31.0.0/20 Default

[Create a subnet](#)

Availability Zone distribution - new

Auto Scaling automatically balances instances across Availability Zones. If launch failures occur in a zone, select a strategy.

☒ **Balanced best effort**
If launches fail in one Availability Zone, Auto Scaling will attempt to launch in another healthy Availability Zone.

☐ **Balanced only**
If launches fail in one Availability Zone, Auto Scaling will continue to attempt to launch in the unhealthy Availability Zone to preserve balanced distribution.

[Cancel](#)

[Skip to review](#)

[Previous](#)

[Next](#)

[EC2](#) > [Auto Scaling groups](#) > Create Auto Scaling group

1

- Step 1
Choose launch template
- Step 2
Choose instance launch options
- Step 3 - optional
Integrate with other services
- Step 4 - optional
Configure group size and scaling
- Step 5 - optional
Add notifications
- Step 6 - optional
Add tags
- Step 7
Review

Integrate with other services - optional [Info](#)

Use a load balancer to distribute network traffic across multiple servers. Enable service-to-service communications with VPC Lattice. Shift resources away from impaired Availability Zones with zonal shift. You can also customize health check replacements and monitoring.

Load balancing [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

Select Load balancing options

☐ **No load balancer**
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☐ **Attach to an existing load balancer**
Choose from your existing load balancers.

☒ **Attach to a new load balancer**
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to a new load balancer

Load balancer type

Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, [visit the Load Balancing console](#).

☒ **Application Load Balancer**
HTTP, HTTPS

☐ **Network Load Balancer**
TCP, UDP, TLS

Load balancer name

Name cannot be changed after the load balancer is created.

SristiAuto-1

Must be unique within the Region. Use 1–32 alphanumeric characters and hyphens. Can't start with "internal-" or begin or end with a hyphen.

Load balancer scheme

Scheme cannot be changed after the load balancer is created.

☐ Internal

☒ Internet-facing

Network mapping

Your new load balancer will be created using the same VPC and Availability Zone selections as your Auto Scaling group. You can select different subnets and add subnets from additional Availability Zones.

Network mapping

Your new load balancer will be created using the same VPC and Availability Zone selections as your Auto Scaling group. You can select different subnets and add subnets from additional Availability Zones.

VPC

vpc-0a874208a064ece04 [Info](#)

Availability Zones and subnets

You must select a single subnet for each Availability Zone enabled. Only public subnets are available for selection to support DNS resolution.

☒ aps1-az3 (ap-south-1b)

Select a subnet
subnet-0f94e0905aad52965

☒ aps1-az2 (ap-south-1c)

Select a subnet
subnet-0e38d5355eecbfc1a

☒ aps1-az1 (ap-south-1a)

Select a subnet
subnet-0e71abe105a1ac286

Listeners and routing

If you require secure listeners, or multiple listeners, you can configure them from the [Load Balancing console](#) after your load balancer is created.

Protocol

HTTP

Port

80

Default routing (forward to)

Create a target group

New target group name

An instance target group with default settings will be created.

SristiAuto-1

Must be unique within the Region. Use 1–32 alphanumeric characters and hyphens. Can't begin or end with a hyphen.

Health checks

Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.

EC2 health checks

[Always enabled](#)

Additional health check types - optional [Info](#)

☒ Turn on Elastic Load Balancing health checks Recommended

Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.

[EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing. To avoid unexpected terminations, first verify the settings of these health checks in the \[Load Balancer console\]\(#\).](#)

☐ Turn on VPC Lattice health checks

VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.

☐ Turn on Amazon EBS health checks

EBS monitors whether an instance's root volume or attached volume stalls. When it reports an unhealthy volume, EC2 Auto Scaling can replace the instance on its next periodic health check.

Health check grace period [Info](#)

This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.

180

seconds

[Cancel](#)

[Skip to review](#)

[Previous](#)

[Next](#)

- Step 1
Choose launch template
- Step 2
Choose instance launch options
- Step 3 - optional
Integrate with other services
- Step 4 - optional
Configure group size and scaling
- Step 5 - optional
Add notifications
- Step 6 - optional
Add tags
- Step 7
Review

Configure group size and scaling - optional [Info](#)

Define your group's desired capacity and scaling limits. You can optionally add automatic scaling to adjust the size of your group.

Group size [Info](#)

Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.

Desired capacity type

Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.

Units (number of instances)

Desired capacity

Specify your group size.

2

Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity

2

Equal or less than desired capacity

Max desired capacity

3

Equal or greater than desired capacity

Automatic scaling - optional

Choose whether to use a target tracking policy [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

☐ No scaling policies

Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.

☒ Target tracking scaling policy

Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling policy name

Target Tracking Policy

Metric type [Info](#)

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

Target value

20

Instance warmup [Info](#)

180

seconds

☐ Disable scale in to create only a scale-out policy

Instance maintenance policy [Info](#)

Control your Auto Scaling group's availability during instance replacement events. This includes health checks, instance refreshes, maximum instance lifetime features and events that happen automatically to keep your group balanced, called rebalancing events.

Choose a replacement behavior depending on your availability requirements

Mixed behavior

☒ No policy

For rebalancing events, new instances will launch before terminating others. For all other events, instances terminate and launch at the same time.

Prioritize availability

☐ Launch before terminating

Launch new instances and wait for them to be ready before terminating others. This allows you to go above your desired capacity by a given percentage and may temporarily increase costs.

Control costs

☐ Terminate and launch

Terminate and launch instances at the same time. This allows you to go below your desired capacity by a given percentage and may temporarily reduce availability.

Flexible

☐ Custom behavior

Set custom values for the minimum and maximum amount of available capacity. This gives you greater flexibility in setting how far below and over your desired capacity EC2 Auto Scaling goes when replacing instances.

- Step 1
● Choose launch template
- Step 2
● Choose instance launch options
- Step 3 - optional
● Integrate with other services
- Step 4 - optional
● Configure group size and scaling
- Step 5 - optional
● **Add notifications**
- Step 6 - optional
● Add tags
- Step 7
● Review

Add notifications - optional [Info](#)

Send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group.

[Add notification](#)[Cancel](#)[Skip to review](#)[Previous](#)[Next](#)

- Step 1
● Choose launch template
- Step 2
● Choose instance launch options
- Step 3 - optional
● Integrate with other services
- Step 4 - optional
● Configure group size and scaling
- Step 5 - optional
● Add notifications
- Step 6 - optional
● **Add tags**
- Step 7
○ Review

Add tags - optional [Info](#)

Add tags to help you search, filter, and track your Auto Scaling group across AWS. You can also choose to automatically add these tags to instances when they are launched.

❗ You can optionally choose to add tags to instances (and their attached EBS volumes) by specifying tags in your launch template. We recommend caution, however, because the tag values for instances from your launch template will be overridden if there are any duplicate keys specified for the Auto Scaling group.

Tags (0)

No tags

[Add tag](#)

50 remaining

[Cancel](#)[Previous](#)[Next](#)

- Step 1
● Choose launch template
- Step 2
● Choose instance launch options
- Step 3 - optional
● Integrate with other services
- Step 4 - optional
● Configure group size and scaling
- Step 5 - optional
● Add notifications
- Step 6 - optional
● Add tags
- Step 7
● **Review**

Review [Info](#)

Step 1: Choose launch template [Edit](#)

Group details

Auto Scaling group name
SristiAuto

Launch template

Launch template
[SristiTemp](#)
lt-0283d209c61392956

Version
Latest

Description
SristiTemp

Step 2: Choose instance launch options [Edit](#)

Network

VPC
[vpc-0a874208a064ece04](#)

Availability Zones and subnets

Availability Zone	Subnet	Subnet CIDR range
aps1-az1 (ap-south-1a)	subnet-0e71abe105a1ac286	172.31.32.0/20
aps1-az2 (ap-south-1c)	subnet-0e38d5355eebcf1a	172.31.16.0/20
aps1-az3 (ap-south-1b)	subnet-0f94e0905aad52965	172.31.0.0/20

Step 6: Add tags [Edit](#)

Tags (0)

Key	Value	Tag new instances
No tags		

[Preview code](#)[Cancel](#)[Previous](#)[Create Auto Scaling group](#)



Auto Scaling groups (1) [Info](#)

Last updated
less than a minute ago

[Launch configurations](#)[Launch templates](#)[Actions](#)[Create Auto Scaling group](#)

< 1 > 

☐	Name	Launch template/configuration	Instances	Status	Desired capacity	Min	Max	Availability Zones	Created
☐	SristiAuto	SristiTemp Version Latest	0	Updating capacity...	2	2	3	3 Availability Zones	Wed A...



```
GNU nano 7.2                               infy.sh *
#!/bin/bash
while(true)
do
    echo "Inside loop"
done
```

Help Write Out Where Is Cut Execute Location M-U Undo
Exit Read File Replace Paste Justify Go To Line M-E Redo

```
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
Inside loop
```

Launch an instance | EC2 | ap-south-1 | Instances | EC2 | ap-south-1 | bk2cridufbe3w04ez3uxsod4u0 | bhowmicksristi29-beep/Repository | + | Ask Gemini

ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#Instances:v=3;\$case=tags:true%5C,client:false;\$regex=tags:false%5C,client:false

aws | Search | [Alt+S] | Asia Pacific (Mumbai) | Sristi Bhowmick (0217-5467-6955) | Sristi Bhowmick

EC2 > Instances

Instances (2/5) Info Last updated 21 minutes ago Connect Instance state Actions Launch instances

Choose filter set

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
☒		i-041b045531c742ed1	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1c	ec2-3-110-126-236.aj
☐	Sristi10	i-0e8237ea5307ef08a	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-13-232-92-252.aj
☐	Sristi12	i-079b6e1ffc94733d9	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-65-0-204-233.ap
☒		i-085020bb3a02bd613	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-3-108-191-143.aj
☐	Sristi14	i-03fcc8f53001c5048	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-13-206-100-176.aj

2 instances selected

☐ Alarm recommendations ☒ Investigate with AI - new 1h 3h 12h 1d 3d 1w Custom UTC timezone Explore related

CPU utilization (%)

Percent

0.415%

0.208%

0%

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

Network in (bytes)

Bytes

5.862k

2.931k

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

Network out (bytes)

Bytes

7.57k

3.785k

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

Network packets in (count)

Count

43.2

21.6

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

Network packets out (count)

Count

41.4

1

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

Metadata no token (count)

Count

1

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

CPU credit usage (count)

Count

0.04

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

CPU credit balance (count)

Count

68.34

0

09:30 10:00 10:30

i-085020bb3a02bd613

i-041b045531c742ed1

CloudShell | Feedback | Console Mobile App | © 2026, Amazon Web Services, Inc. or its affiliates. | Privacy | Terms | Cookie preferences

Launch an instance | EC2 | ap-south-1 | bk2crldufbe3wwo4ez3uxsod4u0 | bhowmicksristi29-beep/Repos | Ask Gemini


ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#Instances:v=3,\$case=tags:true%5C,client:false,\$regex=tags:false%5C,client:false

CPU utilization (%)

5 minutes | Average | 1h 3h 12h 1d 3d 1w Custom | UTC timezone

Percent

0.415%
0.311%
0.208%
0.104%
0%



09:30 09:35 09:40 09:45 09:50 09:55 10:00 10:05 10:10 10:15 10:20 10:25 10:30

i-085020bb3a02bd613 i-041b045531c742ed1

View in metrics | Close

Launch an instance | EC2 | ap-south-1 | bk2crldufbe3wwo4ez3uxsod4u0 | bhowmicksristi29-beep/Repos | Ask Gemini

ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#Instances:v=3,\$case=tags:true%5C,client:false,\$regex=tags:false%5C,client:false

aws

Search [Alt+S]

Asia Pacific (Mumbai) | Sristi Bhowmick (0217-5467-6955) | Sristi Bhowmick

EC2 > Instances

Dashboard
AWS Global View
Events

Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Capacity Manager

Images
AMIs
AMI Catalog

Elastic Block Store
Volumes
Snapshots
Lifecycle Manager

Network & Security
Security Groups
Elastic IPs

Instances (2/5) Info

Last updated 22 minutes ago | Connect | Instance state | Actions | Launch instances

Stop instance
Start instance
Reboot instance
Hibernate instance
Terminate (delete) instance

	Name	Instance ID	Instance state	Instance type	Status check		Availability Zone	Public IPv4 DNS
☒		i-041b045531c742ed1	Running	t3.micro	3/3 checks passed		ap-south-1a	ec2-3-110-126-236.ap-south-1.amazonaws.com
☐	Sristi10	i-0e8237ea5307ef08a	Running	t3.micro	3/3 checks passed		ap-south-1a	ec2-3-110-126-236.ap-south-1.amazonaws.com
☐	Sristi12	i-079b6e1ffc94733d9	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-65-0-204-233.ap-south-1.amazonaws.com
☒		i-085020bb3a02bd613	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-3-108-191-143.ap-south-1.amazonaws.com
☐	Sristi14	i-03fcc8f53001c5048	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1a	ec2-13-206-100-176.ap-south-1.amazonaws.com


2 instances selected

Alarm recommendations | Investigate with AI - new | 1h 3h 12h 1d 3d 1w Custom | UTC timezone | Explore related

CPU utilization (%)

Percent

0.415%
0.208%
0%




09:30 10:00 10:30

i-085020bb3a02bd613 i-041b045531c742ed1

Network in (bytes)

Bytes

5.862k
2.931k




09:30 10:00 10:30

i-085020bb3a02bd613 i-041b045531c742ed1

Network out (bytes)

Bytes

7.57k
3.785k




09:30 10:00 10:30

i-085020bb3a02bd613 i-041b045531c742ed1

Network packets in (count)

Count

43.2
21.6



09:30 10:00 10:30

i-085020bb3a02bd613 i-041b045531c742ed1

Network packets out (count)

Count

41.4



Metadata no token (count)

Count


1



CPU credit usage (count)

Count


0.04



CPU credit balance (count)

Count

68.34



CloudShell | Feedback | Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. | Privacy | Terms | Cookie preferences